

LOAD SELECTION & PROFITABILITY MODEL

Version 3 — cost stack reconciled to owner-operator settlement history

A weighted multi-criteria decision system for freight selection under real operating constraints

Business analysis and decision-support portfolio · Christopher D. Sheppard · Tee 2 Green Golf Labs

Built on eighteen completed loads from my own operation. Cost-per-mile inputs are derived from my own settlement history, using multi-year averages for the fixed and variable components and the most recent three weeks as current-period figures. These are actual operating numbers from a single owner-operator. They are not third-party audited results and, except for the explicitly labeled tire reserve, they are not industry benchmarks.

THE OPERATING PROBLEM

A freight offer arrives as one number. The rate. Almost everything that determines whether the load is worth running sits outside that number: empty miles to reach the shipper, fuel burn at weight, hours surrendered to a slow receiver, and whether the destination has freight moving back out of it. Two loads posted at the same rate per mile can differ by several hundred dollars in what they actually leave behind. The call has to be made in minutes, repeatedly, under fatigue, and the attractive answer and the profitable answer disagree often enough to matter.

I have made this decision a few thousand times across eight years of nationwide flatbed operation. The judgment worked. It was also unwritten, which meant no one could check it, I could not hand it to a driver I was training, and it degraded exactly when I was most tired. This model exists to get that reasoning onto paper where it can be argued with.

THREE VERSIONS, TWO CORRECTIONS FROM EVIDENCE

Version 1 ran on seven representative lanes and illustrative cost figures. It looked correct. That was the flaw. Nothing in it could be shown wrong, because nothing in it had been tested against a decision I had actually made.

Version 2 rebuilt it on 18 completed trips from Prime app history — 19,214 miles and \$41,501.02 — with the score I assigned each load after delivery. That exposed a phantom \$0.280 per mile truck payment on equipment I own outright, a single-point fuel assumption that ignored weight class, and a destination-value weight set far too low.

Version 3 source-checked the cost stack line by line against my own settlement history. Fixed and variable components use multi-year averages, and the current-period figures use the most recent three weeks. The explicitly labeled tire reserve is the one industry placeholder, because tire spend does not appear on the settlement.

What the settlement confirmed

The 18-load sample blends to \$2.1599 per mile. The statement reads \$2.163 year to date across 56,752 miles. A 0.15% difference.

That matters more than it sounds. It means the eighteen loads are a fair slice of the year rather than a flattering selection, which is the first thing a reader should suspect about a sample somebody chose themselves.

WHAT THE SETTLEMENT DISPROVED

- **Maintenance.** Reserved at \$0.180 per mile. Actual is \$0.065 year to date and \$0.070 lifetime. A 2.8x over-reserve carried through two versions of this model without anyone noticing, because it sat inside a total that looked reasonable.
- **Tires — and I got this one wrong twice.** The statement shows a tire fund accrual and a fund dispersal offsetting to the cent, and an earlier version of this model read that as evidence that tires cost nothing. Prime runs a tire fund for lease drivers, not owner-operators — for my unit those lines are vestigial. I buy tires personally out of a working emergency fund, so they never touch the settlement. The statement does not show a zero cost. It shows no cost, which is a different thing, and the model carries \$0.045 per mile as a reserve because of it.
- **Diesel.** Modelled at \$3.70 a gallon. The settlement week shows \$4.933 at the pump against a Prime discount of \$0.857, for a net of \$4.076. Working backwards from the year-to-date fuel line gives a blended 7.20 mpg, which sits exactly where an operator running mostly heavy flatbed with some light freight should land.

- **Annual miles.** Assumed at 120,000 with nothing behind the number. The statement gives a lifetime average of 101,057, and I did not adopt that either. See the section below.

The errors ran in opposite directions and mostly cancelled. The version 2 all-in figure was accurate to within eight cents for entirely the wrong reasons. That is the argument for reconciling a model to source records instead of admiring it, and it is the reason this document has a version 3 rather than a polish of version 2.

THE ONE STATEMENT FIGURE I DID NOT ADOPT

The operating statement reports 1,943.413 miles per week across 237 weeks in service, which annualises to 101,057. That number is source-checked and I still did not use it, because it contains an extended family-leave year in which I was off roughly seven months against a normal two. A lifetime average of a business that shut down for half a year is a real measurement of the past. It is not a run rate, and fixed cost per mile divides by a run rate.

Two reconstructions were run against each other rather than picking one and hoping. Removing that year at roughly 21.7 working weeks leaves 185 weeks averaging 2,167 to 2,196 miles, or 111,000 to 114,000 a year. Building it up independently from 44 working weeks at 2,500 to 2,600 miles gives 110,000 to 114,400. One method from the top, one from the bottom, landing within 3,000 miles of each other.

Annual miles	Fixed CPM	Heavy break-even	What it represents
98,000	\$0.1299	\$0.952	A low-utilization year like the family-leave one
101,057	\$0.1260	\$0.948	Raw lifetime average from the statement
114,400	\$0.1113	\$0.934	Model default: 44 working weeks at 2,600 miles
132,000	\$0.0965	\$0.919	44 weeks at 3,000, the top of the reported range

One caveat carries into this from the section below: the 237 weeks behind these reconstructions are the current tractor only, so the earlier years cannot confirm or contradict them.

The whole spread is \$0.033 per mile, so this assumption is worth about a third of what the maintenance line is worth and the model is not fragile to it. The direction of the risk is worth stating plainly though: a higher annual figure lowers fixed cost per mile, which lowers the bar a load has to clear. Assume a year that does not happen and the truck under-recovers its fixed cost. The driver is exposed on the assumptions tab as working weeks times miles per working week, so it can be argued with rather than trusted.

UNIT ECONOMICS: THE NUMBER THE LOAD MUST BEAT

Cost per mile splits into variable and fixed. Fixed converts to a per-mile figure through a derived 114,400 annual miles, so a low-utilization year correctly raises the floor. Both apply to total miles, loaded and empty, because deadhead consumes fuel, hours, and equipment life while earning nothing. Fuel is calculated per load from the weight class, which produces three break-even figures rather than one.

Component	Heavy	Medium	Light	Source on the operating statement
Variable — fuel, net of discount	\$0.582	\$0.480	\$0.408	\$4.076/gal net at 7.0 / 8.5 / 10.0 mpg
Variable — DEF, net of discount	\$0.023	\$0.023	\$0.023	TRACTOR DEF less DEF DISCOUNT, YTD
Variable — state fuel tax	\$0.007	\$0.007	\$0.007	STATE FUEL TAX, YTD
Variable — maintenance & repair	\$0.120	\$0.120	\$0.120	Actual \$0.065; set higher by choice, see below
Variable — tires (reserve)	\$0.045	\$0.045	\$0.045	NOT a statement line — see below
Variable — tolls, Qualcomm, misc	\$0.045	\$0.045	\$0.045	MISCELLANEOUS less advances and fund transfers
Fixed — truck payment	\$0.000	\$0.000	\$0.000	Lease completed January 2025. RENT PAYMENT \$0.
Fixed — insurance, permits, taxes	\$0.111	\$0.111	\$0.111	Seven statement lines over 114,400 annual miles
ALL-IN COST PER MILE	\$0.934	\$0.831	\$0.759	See the reconciliation immediately below

Maintenance is the one line I did not take at face value. Year to date reads \$0.065 and lifetime reads \$0.070. Those agreeing is weaker evidence than it first looks, because both come from the same tractor over the same period — they confirm each other rather than the underlying rate.

Two things push the real number higher. The truck has 460,589 miles and the lease finished in January 2025, so the warranty window that produced those years is behind it. And I ran this tractor gently the whole time I was buying it, which the one before it did not get. A maintenance figure drawn only from the careful years is a floor, not an average. One turbo erases three years of \$0.065, so the model runs at \$0.120 as a deliberate over-reserve. That is a judgement call rather than a measurement, and it is flagged in the workbook rather than buried in a cell.

Break-even moved from \$1.342 in version 1 to \$0.934 here. That is often reported as though removing the truck payment did all of it. It did not, and the bridge matters:

Change	Effect on break-even	Running total
Version 1 all-in cost per mile	—	\$1.342
Truck payment removed (owned outright)	−\$0.280	\$1.062
Fuel re-priced to actual, net of discount	−\$0.010	\$1.052
Maintenance reserve cut from \$0.180 to \$0.120	−\$0.060	\$0.992
Tolls and securement replaced by DEF, fuel tax, misc	+\$0.020	\$1.012
Fixed cost re-derived from statement lines	−\$0.079	\$0.934
VERSION 3 ALL-IN, HEAVY LOAD	−\$0.408 total	\$0.934

The truck payment is \$0.280 of a \$0.408 move, about 69%. The rest came from reconciliation and revised assumptions. Attributing the whole correction to one line would overstate a single finding and bury four others.

Why the truck payment still matters most

Against a comparable owner-operator carrying a \$33,600 annual note, owning the equipment outright is \$0.2937 per mile of structural advantage — \$5,643 across this sample and \$16,668 across the 56,752 miles run so far this year.

That is a competitive position rather than an accounting fix. Loads a financed operator has to refuse are still profitable for me, which is most of the reason Florida freight is worth taking when other flatbed drivers avoid it over the deadhead out.

HOW THE MODEL TIES TO THE STATEMENT

Two figures answer two different questions, and collapsing them into one would be misleading.

	Model	Statement	Gap
Like-for-like — model set to statement values	\$0.832	\$0.838	\$0.006
As configured — the model I actually run	\$0.917	\$0.838	\$0.085

The first line is the arithmetic check. Set maintenance to the actual \$0.065, tires to the zero the statement shows, and divide fixed cost by the statement's own 101,057 annual miles, and the model returns \$0.832 per mile at the year-to-date blended 7.20 mpg. The statement's operating cost over the same scope, less advances and emergency-fund transfers that are not operating costs, is \$0.838. Six tenths of a cent apart. The formulas tie.

The second line is what I actually run, and it sits \$0.085 above the statement. Every cent of that is an itemized choice, not a discrepancy: maintenance reserved \$0.055 above actual, a \$0.045 tire reserve the statement does not carry, and \$0.015 back from normalizing annual miles. Quoting the \$0.006 on its own would imply the model matches the statement. It does not, and it should not — two of those three adjustments exist precisely because the statement understates what the truck costs to run.

RECONCILING ONE LOAD END TO END

One trip in the sample appears on this settlement, so it can be reconciled line by line against what the app displayed when I accepted it. It is the only load of the eighteen that can be checked, and it is worth showing in full because the discrepancies are not where I expected them.

Line	App estimate	Settlement	Delta	Note
Loaded miles	894	894	exact	Matches to the mile.
Empty miles	265	204	-61	App showed more. Cause unconfirmed.
Total miles	1,159	1,098	-61	5.3% fewer on this trip. Not a pattern.
Linehaul (owner share)	\$2,367.83	\$2,367.94	+\$0.11	72% split of \$3,288.81. Rounding only.
Fuel surcharge	\$439.20	\$439.20	—	Exact.
Stop / pickup pay	\$144.00	\$144.00	—	Exact.
Longevity pay	—	\$21.96	+\$21.96	Does not appear in the app.
Toll reimbursement	—	\$13.33	+\$13.33	Does not appear in the app.
TOTAL REVENUE	\$2,951.03	\$2,986.43	+\$35.40	App understates by the two lines above.
Rate per total mile	\$2.546	\$2.720	+\$0.174	On this trip. One observation.

The revenue side is close enough to trust: the linehaul matches to eleven cents, and the two missing lines are pay types the app simply does not display. The mileage side is not. Loaded miles matched exactly while empty miles came in 61 short, which rules out a rounding artifact and leaves either a repower-specific quirk or a different mileage basis between the app and the settlement.

I have not resolved which, so the other seventeen keep the mileage the app reported and the discrepancy is disclosed rather than quietly corrected. One reconciled trip is one observation. It does not establish that the app runs high generally, and no correction has been applied to the rest on the strength of it. Whether this is a repower artifact or a systematic difference in mileage basis is an open question that seventeen more settlements would answer.

SCORING MODEL

Eight dimensions weighted to 100%. Three are calculated from the load economics. Five are operator judgment scored 1 to 5 at the moment of the offer, because reload strength, facility behavior, securement burden, road conditions, and schedule feasibility are field knowledge that appears on no rate sheet. Keeping calculated and judged visibly separate shows how much of the decision is arithmetic and how much is experience.

Dimension	Weight	Type	Score 5 / Score 1 boundary
Rate quality — net \$ per total mile	25%	Calculated	>= \$1.60 / < \$1.05
Destination value — reload, hub, or home	25%	Judgment	Hub or home / Dead market
Time efficiency — net \$ per hour committed	15%	Calculated	>= \$78 / < \$50
Deadhead efficiency — empty as % of total	10%	Calculated	<= 5% / > 28%
Dwell and detention risk	10%	Judgment	Drop & hook / Known problem facility
Securement and handling burden	5%	Judgment	No tarp / Oversize with escort
Route conditions	5%	Judgment	Flat, cheap fuel / Grades, no parking
Hours-of-service and appointment feasibility	5%	Judgment	Ample margin / Not legally possible

The rate band has been re-cut in every version, which is the point rather than an embarrassment. Version 1 called \$0.60 a five against a \$1.342 cost floor. Version 2 raised it to \$1.50 when the phantom truck payment came out. Version 3 raises it again to \$1.60 because the settlement-derived stack came in lower still. Left at the version 1 setting, seventeen of eighteen loads score a five and the dimension stops measuring anything.

THE THREE OVERRIDE GATES

A weighted average can be dragged upward by strong dimensions while the load remains unacceptable. Three conditions override the score outright and force a decline. This is the difference between a scoring exercise and a decision system.

- **Gate 1 — negative net contribution.** The load does not cover its own all-in cost. No amount of reload strength repairs a negative contribution.
- **Gate 2 — deadhead above tolerance.** Set at 35% of total miles. Empty miles are unbillable and consume hours of service that cannot be recovered.
- **Gate 3 — infeasible or blacklisted.** If the appointment cannot legally be made the rate is irrelevant, and accepting creates compliance exposure. A blacklisted facility trips the same gate. Three sit on that list, each for a documented incident, and they are permanent declines regardless of what they pay.

THE SAMPLE: 18 COMPLETED TRIPS

Net dollars per mile is after the settlement-derived all-in cost. Model score is on the 1 to 5 weighted scale; operator score is the 1 to 10 rating I assigned each load after delivery.

Load	Lane	Gross \$/mi	Net \$/mi	DH %	Model	Op	Decision
L-01	Visalia CA → Peoria AZ	\$2.08	\$1.25	27.5%	3.50	8	NEGOTIATE
L-02	Tempe AZ → Rayville LA	\$2.15	\$1.39	1.8%	4.05	9	ACCEPT
L-03	Natchitoches LA → Salt Lake City UT	\$2.36	\$1.43	8.1%	4.40	8	ACCEPT
L-04	Cedar City UT → Brentwood CA	\$2.52	\$1.58	26.9%	4.40	10	ACCEPT
L-05	Oakland CA → Mead CO	\$2.40	\$1.46	3.7%	4.05	9	ACCEPT
L-06	Pueblo CO → Stillwater OK	\$1.42	\$0.49	22.1%	2.40	5	DECLINE — score
L-07	Maize KS → Sikeston MO	\$2.01	\$1.07	20.9%	2.85	9	DECLINE — score
L-08	Troy IL → Sulphur LA	\$1.99	\$1.23	18.2%	3.60	10	NEGOTIATE
L-09	Joyce LA → Elmendorf TX	\$2.37	\$1.44	27.3%	3.80	8	ACCEPT
L-10	Houston TX → Ontario CA	\$2.03	\$1.10	12.0%	2.75	7	DECLINE — score
L-11	Visalia CA → Roseburg OR	\$1.80	\$1.04	26.3%	2.35	8	DECLINE — score
L-12	Riddle OR → Midlothian TX	\$2.00	\$1.06	1.1%	3.15	5	NEGOTIATE
L-13	Terrell TX → Fort Bliss TX	\$2.93	\$2.00	7.7%	4.65	10	ACCEPT
L-14	Socorro TX → Apollo Beach FL	\$1.89	\$0.96	1.0%	3.10	10	NEGOTIATE
L-15	Lake City FL → Charlottesville VA	\$2.29	\$1.53	21.3%	3.00	9	NEGOTIATE
L-16	Winchester VA → Pittston PA	\$2.94	\$2.18	31.2%	4.50	10	ACCEPT
L-17	Reading PA → Springfield MO	\$1.89	\$1.06	7.7%	3.55	9	NEGOTIATE
L-18	Blytheville AR → W Palm Beach FL	\$2.55	\$1.61	22.9%	4.05	9	ACCEPT
BLENDED / 19,214 mi • \$41,501.02		\$2.16	\$1.28	12.7%	3.56	8.50	

Total net contribution across the sample is \$24,504 against \$41,501.02 of revenue. Deadhead is 12.7%.

TESTING AGAINST MY OWN RATINGS, AND THE RESULT I DID NOT WANT

First, what this test is and is not. All 18 loads were accepted, and I scored each one after delivery. This compares model output against retrospective quality ratings. It is not a test against accept-or-decline decisions, and nothing in it shows whether the model would have caught a load worth refusing.

Correlation between the model score and my rating across the 18 loads is 0.56. Squared, the model accounts for roughly a third of the variance in those ratings.

Net dollars per mile, on its own, correlates at 0.62.

A single arithmetic column tracks my judgment slightly better than the eight-dimension weighted model built on top of it. I could close that gap by re-tuning the weights against these eighteen rows and the number would improve. It would also stop meaning anything, because at $n=18$ that is fitting noise and calling it insight. The weights stay where reasoning put them, and the result gets reported as it came out.

Measure	Value	Reading
Rating vs. gross \$/mile	0.57	Rate alone, before costing the load
Rating vs. net \$/mile	0.62	Costing it properly helps
Rating vs. full model score	0.56	Eight dimensions, no better than net rate
Variance in ratings explained	30.9%	Two thirds of the rating sits outside the model
Lowest four — my rating	L-06, L-12, L-10, L-01	Two of four overlap with the model
Lowest four — model	L-11, L-06, L-10, L-07	
Highest four — my rating	L-04, L-08, L-13, L-14	Two of four overlap with the model
Highest four — model	L-13, L-16, L-03, L-04	

The disagreements are more useful than the overlaps. L-14 and L-08 both rated 10 and both land in the model's bottom half. L-14 delivered to Apollo Beach and put me in my own house for three days. L-08 was 2,500 pounds of strut that ran 11 mpg and cleared the receiver in under thirty minutes. L-07 the model declines outright; three drivers had already refused it when it came to me, and taking it bought discretion with my fleet manager that is worth more over a year than the margin I gave up on one trip.

Two of those three misses are structural and should stay that way. A weighted score built on load economics cannot see the value of being home on a Saturday, and a model that claims to price that is lying about what it knows. They are handled outside the score instead: a weekly revenue floor, and a written exception for loads taken to bank goodwill.

Why the model survives a result like this

It was never going to beat my own ratings on loads I have already run. It does three things a rating cannot. It is reviewable, so any decision can be reconstructed months later. It is teachable, so the reasoning transfers to a driver who has not spent eight years acquiring it. And it is stable at 2 a.m. in a way that I am not.

The negative result is also the reason the artifact is worth reading. A results section reporting 0.94 with no threats to validity would tell you the model was tuned until it agreed with itself.

WHAT THE MODEL CHANGES

- **It prices time.** Net dollars per hour committed carries 15%, which penalizes a well-paying load that will burn eleven hours at a slow facility. L-06 clears its costs and still scores 2.40 because it returns \$23 an hour against a sample average of \$61.
- **It makes deadhead a cost, not a rule of thumb.** Full cost per mile applies to empty miles. L-16 carries 31.2% deadhead and still scores 4.50, because 330 total miles at \$2.94 gross absorbs it. The model gets that trade right where a flat deadhead rule would not.
- **It prices weight through commodity.** Foam board at 5,000 pounds returns 10 mpg against 7.0 for a loaded flatbed. At \$4.076 net diesel that is \$0.17 per mile of margin a single-MPG model hands to the wrong load.
- **It separates profitable from acceptable.** L-12 clears every gate, returns \$2,194 net, and still only earns a counter. Grades through the Sierra Nevada, no signal, and nowhere legal to shut down are real costs that never reach the rate confirmation.
- **It is teachable.** The rules are written down, so the reasoning can be handed to a new driver instead of acquired across three years of expensive mistakes.

ENTERPRISE TRANSLATION

The freight context is incidental. What the artifact demonstrates is a repeatable pattern for high-frequency decisions made under incomplete information and time pressure: establish the true unit cost, weight the criteria that actually drive the outcome, keep calculated separate from judged, put hard gates around the failures no weighted average should ever approve, then test it against whatever record of past outcomes exists and publish what the test said.

The version 3 step matters most here. Anyone can build a cost model from published benchmarks. Reconciling one to settlement-derived operating history, finding four material errors in your own work, and rewriting rather than defending is the part that transfers.

Freight decision	Enterprise AI platform evaluation
All-in cost per mile	Total cost of ownership — licensing, integration, support, retraining
Reconciling the cost stack to settlement history	Reconciling a TCO model to actual invoices and usage telemetry
Rate quality	Measured task quality and business outcome, not demo performance
Time efficiency	Time to value; engineering hours consumed per unit of benefit
Deadhead percentage	Unrecoverable overhead — rework, migration, abandoned pilots
Destination value	Platform roadmap, vendor viability, what the next decision costs
Dwell and detention risk	Support responsiveness, incident history, escalation paths
Securement burden	Configuration load, identity and access work, change management
Route conditions	Organizational readiness, competing initiatives, adoption friction
Gate 1 — negative contribution	Fails the business case; no measurable outcome to point at
Gate 2 — deadhead over tolerance	Integration or data-residency requirements exceed tolerance
Gate 3 — infeasible or blacklisted	Fails security, privacy, or regulatory review. Not negotiable.

WHAT THE SOURCE DATA COVERS

Every lifetime figure quoted here means life of the current tractor, not life of my time at this carrier, and the distinction is large enough that leaving it implicit would be misleading.

Period	Status in these figures
March 2018 — July 2018	In a trainer's truck, running the 50,000 miles required to earn an A-seat. Never my miles.
July 2018 — December 2021	My own tractor, roughly three and a half years. Absent from every figure in this document.
January 2022 — present	The current tractor. 237 weeks, 460,589 miles. This is the entire evidence base.

Prime resets lifetime totals when a driver takes a new lease or lease purchase, so the earlier operating history is not recoverable from this settlement. Tenure at the carrier is roughly eight and a half years; the source-checked window is four and a half of them. The prior tractor was turned in showing about 595,000 odometer miles from factory, and how many of those I put on it is not established by anything I can produce, so no rate is derived from that number here.

The gap is not neutral, which is the part worth saying out loud. I drove the previous truck into the ground and treated this one carefully because I was buying it. A cost stack built only from the careful years understates what a truck costs to run. That is the second reason the maintenance line sits above its observed value, and it is a reason a reader would have no way to discover from the settlement alone.

EVIDENCE BOUNDARY

Cost-per-mile inputs in the unit economics table are derived from my own settlement history, using multi-year averages for fixed and variable components and the most recent three weeks as current-period figures. The explicitly labeled tire reserve is the exception: it is an industry placeholder because tire spend does not appear on

the settlement. The source document is not distributed with this case study because it carries bank routing detail, customer names, and a second driver's name.

Revenue and mileage for the 18 loads come from Prime app trip history. One of the eighteen has been reconciled to a settlement; the other seventeen have not, and that single reconciled load differed by 5.3% on total miles and \$35.40 on revenue. One observation cannot establish a pattern and none has been assumed. Ratings are self-assigned after delivery and subjective by construction — outcome leaks into them. L-12 I marked a 5, then revised to a 7 once I separated a \$220 citation from the quality of the decision to accept.

Drive hours derive from an assumed 52 mph average including fuel and inspection stops, not from ELD records. Dwell hours come from recall. The rate and time band boundaries were cut from the observed distribution of these same 18 loads, so the reported correlation is in-sample and optimistic by construction.

The largest gap is unchanged: all 18 loads were accepted. The model has never been scored against an offer I turned down, so nothing here speaks to its behavior on the decline side. Logging refused offers is the next step and the one that would change the most.

One cost line has no source document behind it. The \$0.045 tire reserve is an industry placeholder, carried because Prime runs no tire fund for owner-operators and my tire spend never reaches a settlement. It is the weakest number in the stack.

Every assumption above sits in a labeled cell in the workbook. Replacing any of them recalculates every downstream result.

ABOUT

Christopher D. Sheppard is a Marine Corps veteran, owner-operator, Microsoft PL-900 certified professional, and founder of Tee 2 Green Golf Labs. He builds applied AI systems, decision-support tools, and operating documentation, and is moving from regulated field operations into enterprise AI platform delivery and business systems work.

[linkedin.com/in/christophersheppard](https://www.linkedin.com/in/christophersheppard)